

CITIZEN AI:INTELLIGENT CITIZEN ENAGEMENT PLATFORM

Project Documentation

1. Introduction :

Project title : Citizen Ai:Intelligent Citizen Enagement Platform.

- Team member : MOHAN RAJ G
- Team member : THARUN R
- Team member : DHARMESH N
- Team member : EBIANZER V

2. Project Overview:

- **Purpose:** The goal of the Citizen AI Project is to empower city residents by leveraging AI and real-time data to create a more eco-conscious and connected urban environment. It helps optimize resources like energy, water, and waste, and provides personalized eco-tips to encourage sustainable behaviour among citizens. For city officials, the project serves as a decision-making tool by providing insights, forecasting capabilities, and summaries of complex policies. The project aims to connect technology, governance, and community to build more efficient, resilient, and greener cities.
- **Features:**
 - **Conversational Interface:** This allows for natural language interaction, enabling citizens and officials to ask questions and receive guidance.
 - **Policy Summarization:** Converts long government documents into clear, actionable summaries for easier understanding.
 - **Resource Forecasting:** Uses historical and real-time data to predict future usage of energy, water, and waste.
 - **Eco-Tip Generator:** Recommends daily actions to help users reduce their environmental impact based on their behavior.
 - **Citizen Feedback Loop:** Gathers and analyzes public input to assist with city planning and service enhancements.

- **KPI Forecasting:** Projects key performance indicators to help officials monitor progress and plan strategically.
- **Anomaly Detection:** Acts as an early warning system by identifying unusual patterns in sensor or usage data to flag potential issues.
- **Multimodal Input Support:** Can handle different data types, including text, PDFs, and CSVs, for analysis and forecasting.
- **User-friendly Interface:** An intuitive dashboard built with Streamlit or Gradio UI that allows both citizens and city officials to easily interact with the assistant.

3. Architecture:

- **Frontend (Streamlit):** The frontend is an interactive web UI with multiple pages for dashboards, file uploads, a chat interface, feedback forms, and report viewers. It uses the Streamlit-option-menu library for sidebar navigation, and each page is modularized for scalability.
- **Backend (FastAPI):** This serves as the REST framework for API endpoints that handle document processing, chat, eco-tip generation, and more. It is optimized for asynchronous performance and easy Swagger integration.
- **LLM Integration (IBM Watsonx Granite):** The project uses Granite LLM models from IBM Watsonx for natural language understanding and generation. Prompts are specifically designed to produce summaries, reports, and sustainability tips.
- **Vector Search (Pinecone):** Uploaded policy documents are converted into embeddings using Sentence Transformers and stored in Pinecone. Semantic search is enabled via cosine similarity, letting users search documents using natural language queries.
- **ML Modules (Forecasting and Anomaly Detection):** Lightweight ML models from Scikit-learn are used for forecasting and anomaly detection. Time-series data is parsed, modeled, and visualized using pandas and matplotlib.

4. Setup Instructions:

- **Prerequisites:**

- Python 3.9 or later
- pip and virtual environment tools
- API keys for IBM Watsonx and Pinecone
- Internet access for cloud services
- **Installation Process:**
 - Clone the repository.
 - Install dependencies from requirements.txt.
 - Create and configure a .env file with credentials.
 - Run the backend server using FastAPI.
 - Launch the frontend via Streamlit.
 - Upload data and interact with the modules.

5. Folder Structure:

- app/ - Contains all FastAPI backend logic, including routers, models, and integration modules.
- app/api/ - Subdirectory for modular API routes like chat, feedback, and document vectorization.
- ui/ - Contains frontend components for Streamlit pages and form UIs.
- smart_dashboard.py - The entry script for the main Streamlit dashboard.
- granite_llm.py - Handles all communication with the IBM Watsonx Granite model.
- document_embedder.py - Converts documents to embeddings and stores them in Pinecone.
- kpi_file_forecaster.py - Forecasts future trends for energy/water using regression.

- `anomaly_file_checker.py` - Flags unusual values in uploaded KPI data.
- `report_generator.py` - Constructs AI-generated sustainability reports.

6. Running the Application:

- To start the project, launch the FastAPI server and then run the Streamlit dashboard.
- Navigate through the pages using the sidebar.
- Users can upload documents or CSVs, interact with the chat assistant, and view outputs like reports, summaries, and predictions.
- All interactions are real-time, with the frontend dynamically updating via backend APIs.

7. API Documentation:

- The backend APIs include:
 - `POST /chat/ask` - Accepts a user query and returns an AI-generated message.
 - `POST /upload-doc` - Uploads and embeds documents in Pinecone.
 - `GET /search-docs` - Returns semantically similar policies to a user query.
 - `GET /get-eco-tips` - Provides sustainability tips on selected topics.
 - `POST /submit-feedback` - Stores citizen feedback.
- Each endpoint is documented and tested in Swagger UI.

8. Authentication:

- For demonstration purposes, this version of the project runs in an open environment.
- Secure deployments can include:
 - Token-based authentication (JWT or API keys).

- OAuth2 with IBM Cloud credentials.
 - Role-based access for different user types (admin, citizen, researcher).
- Future enhancements will include user sessions and history tracking.
-
-

•

9. User Interface:

- The interface is minimalist and designed for accessibility for non-technical users.
- Key elements include:
 - A sidebar for navigation.
 - KPI visualizations with summary cards.
 - Tabbed layouts for chat, eco tips, and forecasting.
 - Real-time form handling.
 - PDF report download capability.

10. Testing:

- Testing was conducted in several phases:
 - **Unit Testing:** For prompt engineering functions and utility scripts.
 - **API Testing:** Done via Swagger UI, Postman, and test scripts.
 - **Manual Testing:** To validate file uploads, chat responses, and output consistency.
 - **Edge Case Handling:** To address malformed inputs, large files, and invalid API keys.
- Each function was validated to ensure reliability in both offline and API-connected modes.

11.Source Code Screenshots:

12.Source Output:

13. Future Enhancements:

- **User Sessions and History Tracking:** The project plans to add the ability to track user sessions and interaction history. This will allow for a more personalized experience.
- **Security:** For secure deployments, the project can integrate token-based authentication (JWT or API keys), OAuth2 with IBM Cloud credentials, and role-based access for different users (e.g., admin, citizen, researcher).

